

Alessio Miolla

June 19, 2021

1 Problem description

Q3: Modelling Assessment:

For this task use the Bank Marketing dataset available on the following address:
<http://archive.ics.uci.edu/ml/datasets/Bank+Marketing>

Build 2 binary classification models using any 2 of the following methods (in R or Python): 1. Logistic Regression 2. Random Forest 3. GBM 4. Xgboost 5. Neural Network

Try and minimize overfitting. Compare the performance of both models using ROC graphs, AUCs, confusion matrices. Provide also the full source code and description of any variable transformations or balancing performed.

For the current task I used the bank-additional.csv with 10% of the examples (4119). Logistic Regression and Random Forrest were used as ML models.

2 Import

```
[22]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV, \
    ↳StratifiedKFold, RepeatedStratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
import category_encoders as ce
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.metrics import classification_report
import scikitplot as skplt
from sklearn import datasets, metrics, model_selection, svm
from sklearn.metrics import auc, roc_curve, roc_auc_score, \
    ↳precision_recall_curve, precision_score, f1_score, recall_score, \
    ↳confusion_matrix, accuracy_score
```

3 Load

```
[4]: path_data = 'C:/Users/alessio/Desktop/[REDACTED]/'
df = pd.read_csv(path_data + 'bank-additional.csv', sep=';')
```

```
[5]: df.head()
```

```
[5]:
```

	age	job	marital	education	default	housing	loan	\
0	30	blue-collar	married	basic.9y	no	yes	no	
1	39	services	single	high.school	no	no	no	
2	25	services	married	high.school	no	yes	no	
3	38	services	married	basic.9y	no	unknown	unknown	
4	47	admin.	married	university.degree	no	yes	no	

	contact	month	day_of_week	...	campaign	pdays	previous	poutcome	\
0	cellular	may	fri	...	2	999	0	nonexistent	
1	telephone	may	fri	...	4	999	0	nonexistent	
2	telephone	jun	wed	...	1	999	0	nonexistent	
3	telephone	jun	fri	...	3	999	0	nonexistent	
4	cellular	nov	mon	...	1	999	0	nonexistent	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	-1.8	92.893	-46.2	1.313	5099.1	no
1	1.1	93.994	-36.4	4.855	5191.0	no
2	1.4	94.465	-41.8	4.962	5228.1	no
3	1.4	94.465	-41.8	4.959	5228.1	no
4	-0.1	93.200	-42.0	4.191	5195.8	no

```
[5 rows x 21 columns]
```

4 Exploratory analysis

```
[20]: df.shape #check shape
```

```
[20]: (4119, 21)
```

```
[24]: df['y'].value_counts() #check proportion of the target classes
```

```
[24]: no      3668
      yes      451
      Name: y, dtype: int64
```

```
[39]: #check value counts for each feature
      for x in df.columns:
          print(df[x].value_counts())
```

```
[71]: #check dtype for each feature
for x in df.columns:
    print(df[x].dtype)
```

5 Conversion of Categorical variables

Binary transformation The following variables can be converted into dichotomic variables to optimize the learning of the models and optimize the value of the variables.

```
[6]: #convert categorical in numerical values of the target variable
df['y'] = df['y'].map({'no':0, 'yes':1}).astype('uint8')
```

```
[7]: #convert binary categorical variables into numerical
df['contact'] = df['contact'].map({'cellular': 1, 'telephone': 0}).
    →astype('uint8')
df['loan'] = df['loan'].map({'yes': 1, 'unknown': 0, 'no' : 0}).astype('uint8')
df['housing'] = df['housing'].map({'yes': 1, 'unknown': 0, 'no' : 0}).
    →astype('uint8')
df['default'] = df['default'].map({'no': 1, 'unknown': 0, 'yes': 0}).
    →astype('uint8')
```

```
[9]: #convert previous col into binary col (0 if not contacted, 1 if contacted)
df['previous'] = df['previous'].apply(lambda x: 1 if x > 0 else 0).
    →astype('uint8')
```

```
[10]: #convert poutcome col into binary (1 if succeeded, 0 if not)
df['poutcome'] = df['poutcome'].map({'nonexistent':0, 'failure':0, 'success':1}).
    →astype('uint8')
```

```
[8]: #convert 999 pdays in 0
df['pdays'] = df['pdays'].replace(999, 0)
```

OHE encoding One Hot encoding was used to convert the following categorical variables into dummy variables.

```
[11]: #convert job, week, month into one hot encoding

# function to One Hot Encoding
def encode(data, col):
    return pd.concat([data, pd.get_dummies(col, prefix=col.name)], axis=1)

df = encode(df, df.job)
df = encode(df, df.month)
df = encode(df, df.day_of_week)

df.drop(['job', 'month', 'day_of_week'], axis=1, inplace=True)
```

```
df.drop_duplicates(inplace=True)
```

Ordinal encoding The education level was converted into ordinal encoding to rank the different level of educations

```
[12]: #convert categorical into ordinal encoding
df['education'] = df['education'].map({'illiterate':0, 'unknown':0, 'basic.4y':
    →1, 'basic.6y':1, 'professional.course':2, 'basic.9y':2,
    'high.school':3, 'university.degree':4 }).
    →astype('uint8')
```

Target encoding There is no ordinal relationship between the marital metrics (e.g., single or divorced). In addition, all the metrics are homogenous within the sample. A target encoding was used to optimize the encoding of the variable.

```
[13]: # target variable for testing
y = df.y
# convert marital variable into target encoding --> final training_set
target_encode = ce.target_encoder.TargetEncoder(cols=('marital')).fit(df, y)
training_set = target_encode.transform(df)
# drop target variable to split features from target variable
training_set.drop('y', axis=1, inplace=True)
```

C:\Users\aleccio\anaconda3\envs\wargaming\lib\site-packages\category_encoders\utils.py:21: FutureWarning: is_categorical is deprecated and will be removed in a future version. Use is_categorical_dtype instead

```
elif pd.api.types.is_categorical(cols):
```

```
[14]: training_set.head()
```

```
[14]:
```

	age	marital	education	default	housing	loan	contact	duration	\
0	30	0.100438	2	1	1	0	1	487	
1	39	0.134432	3	1	0	0	0	346	
2	25	0.100438	3	1	1	0	0	227	
3	38	0.100438	2	1	0	0	0	17	
4	47	0.100438	4	1	1	0	1	58	

	campaign	pdays	...	month_mar	month_may	month_nov	month_oct	\
0	2	0	...	0	1	0	0	
1	4	0	...	0	1	0	0	
2	1	0	...	0	0	0	0	
3	3	0	...	0	0	0	0	
4	1	0	...	0	0	1	0	

	month_sep	day_of_week_fri	day_of_week_mon	day_of_week_thu	\
--	-----------	-----------------	-----------------	-----------------	---

0	0	1	0	0
1	0	1	0	0
2	0	0	0	0
3	0	1	0	0
4	0	0	1	0

	day_of_week_tue	day_of_week_wed
0	0	0
1	0	0
2	0	1
3	0	0
4	0	0

[5 rows x 44 columns]

6 ML models

The machine learning models used in the current are Logistic Regression and Random Forrest, respectively. F1 score was adopted as the main criteria since the unbalanced between the two target variables. Indeed, the highly skewed class distribution (89% of class0 versus 11% of class1) can lead the classifier to get a low misclassification rate simply by choosing the majority class (i.e., class 0). Consequently, I decided to get the classifiers with high F1 scores in both classes instead of other criteria such as accuracy, precision, and recall of the majority class0.

6.1 Basic approach

Logistic Regression

```
[150]: #split training and test set
X_train, X_test, y_train, y_test = train_test_split(training_set, y, test_size=0.
→2, random_state=42)
Counter(y_test)
```

```
[150]: Counter({1: 92, 0: 732})
```

```
[151]: sc = StandardScaler()
X_train_transformed = sc.fit_transform(X_train)
X_test_transformed = sc.transform(X_test)
```

```
[153]: #LR
model = LogisticRegression()

#Tuning and finding the best hyperparams
solvers = ['newton-cg', 'lbfgs', 'liblinear']
penalty = ['l2']
c_values = [0.3, 0.6, 0.7]
```

```

#grid search
grid = dict(solver=solvers,penalty=penalty,C=c_values)
grid_search_logistic = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1,
→cv=5, scoring='f1',error_score=0)

#training
clf = grid_search_logistic.fit(X_train_transformed, y_train)
y_pred_train = clf.predict(X_train_transformed)
roc=roc_auc_score(y_train, y_pred_train)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_train, y_pred_train))
print(classification_report(y_train, y_pred_train))

```

roc 0.7324928085128991

prediction_matrix

[[2870 66]

[184 175]]

	precision	recall	f1-score	support
0	0.94	0.98	0.96	2936
1	0.73	0.49	0.58	359
accuracy			0.92	3295
macro avg	0.83	0.73	0.77	3295
weighted avg	0.92	0.92	0.92	3295

[89]:

```

#test
y_test_pred = clf.predict(X_test_transformed)

#performance
roc=roc_auc_score(y_test, y_test_pred)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_test, y_test_pred))
print(classification_report(y_test, y_test_pred))

```

roc 0.6772095509622238

prediction_matrix

[[705 27]

[56 36]]

	precision	recall	f1-score	support
0	0.93	0.96	0.94	732
1	0.57	0.39	0.46	92
accuracy			0.90	824
macro avg	0.75	0.68	0.70	824

weighted avg	0.89	0.90	0.89	824
--------------	------	------	------	-----

The test results show slight overfitting of the model with a global F1 score of 0.7 and AUC of 0.68. The minority class negatively affects the overall performance of the model as highlighted in its precision and recall.

Random Forrest

```
[156]: #split training and test set
X_train, X_test, y_train, y_test = train_test_split(training_set, y, test_size=0.
→2, random_state=42)
Counter(y_test)
```

```
[156]: Counter({1: 92, 0: 732})
```

```
[207]: #RF
random_forest_clf = RandomForestClassifier()

#Tuning and finding the best hyperparams

param_grid = {
    'max_leaf_nodes' : [10, 20, 30],
    'max_depth': [10, 20, 30],
}

#grid search
grid_search_rf = GridSearchCV(estimator = random_forest_clf, param_grid =_
→param_grid, cv=5, scoring='f1', verbose=0, n_jobs=-1)

#training
clf_rf = grid_search_rf.fit(X_train, y_train)
y_pred_train = clf_rf.predict(X_train)
roc=roc_auc_score(y_train, y_pred_train)
print('prediction_matrix', '\n', confusion_matrix(y_train, y_pred_train))
print('roc', roc)
print(classification_report(y_train, y_pred_train))
```

```
prediction_matrix
[[2931   5]
 [ 249 110]]
roc 0.6523518439807822
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	2936
1	0.96	0.31	0.46	359
accuracy			0.92	3295
macro avg	0.94	0.65	0.71	3295

weighted avg	0.93	0.92	0.90	3295
--------------	------	------	------	------

```
[208]: #test
y_pred = clf_rf.predict(X_test)

#performance
roc=roc_auc_score(y_test, y_pred)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
roc 0.5930446661914944
prediction_matrix
[[725   7]
 [ 74  18]]
```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	732
1	0.72	0.20	0.31	92
accuracy			0.90	824
macro avg	0.81	0.59	0.63	824
weighted avg	0.89	0.90	0.88	824

RF model includes different limitations. First of all, there is slight overfitting of the training set that affects the test performance. Moreover, the model is severely biased by the majority class, as showed either by the low number of predictions of class1 and the higher precision/recall of the majority class. More restrictions should be adopted in the model to prevent overfitting, with the risk of decreasing the performance of the model or, conversely, causing underfitting of the model.

6.2 Oversampling

The imbalance between the two target classes led to having poor performance in the minority class. In the current section, the oversampling of the minority class was used to addressing this problem. The Synthetic Minority Oversampling Technique (SMOTE) is a type of data augmentation that duplicates examples in the minority class, although these examples don't add any new information to the model. Doing so, I tested if the previous low predictions on class1 are not due to the relatively little amount of information given in the training set.

Logistic Regression

```
[239]: # summarize class distribution
counter = Counter(y)
print(counter)
#split_training and test set
X_train, X_test, y_train, y_test = train_test_split(training_set, y, test_size=0.
→2, random_state=42)
```



```

#compromise oversampling
over = SMOTE(sampling_strategy=0.5)
X_train, y_train = over.fit_resample(X_train, y_train)
# summarize the new class distribution
counter = Counter(y_train)
print(counter)

```

```

Counter({0: 3668, 1: 451})
Counter({0: 2936, 1: 1468})

```

```

[240]: sc = StandardScaler()
X_train_transformed = sc.fit_transform(X_train)
X_test_transformed = sc.transform(X_test)

```

```

[241]: #LR
model = LogisticRegression()

#Tuning and finding the best hyperparams
solvers = ['newton-cg', 'lbfgs', 'liblinear']
penalty = ['l2']
c_values = [100, 10, 1.0, 0.1, 0.01]

#grid search
grid = dict(solver=solvers,penalty=penalty,C=c_values)
grid_search_logistic = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1,
    ↪cv=5, scoring='f1',error_score=0)

#training
clf = grid_search_logistic.fit(X_train_transformed, y_train)
y_pred_train = clf.predict(X_train_transformed)
roc=roc_auc_score(y_train, y_pred_train)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_train, y_pred_train))
print(classification_report(y_train, y_pred_train))

```

```
roc 0.9189373297002724
```

```
prediction_matrix
```

```
[[2846  90]
```

```
[ 193 1275]]
```

	precision	recall	f1-score	support
0	0.94	0.97	0.95	2936
1	0.93	0.87	0.90	1468
accuracy			0.94	4404
macro avg	0.94	0.92	0.93	4404
weighted avg	0.94	0.94	0.94	4404

```
[242]: #test
y_test_pred = clf.predict(X_test_transformed)

#performance
roc=roc_auc_score(y_test, y_test_pred)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_test, y_test_pred))
print(classification_report(y_test, y_test_pred))
```

```
roc 0.6996020432406747
prediction_matrix
[[698  34]
 [ 51  41]]
```

		precision	recall	f1-score	support
	0	0.93	0.95	0.94	732
	1	0.55	0.45	0.49	92
	accuracy			0.90	824
	macro avg	0.74	0.70	0.72	824
	weighted avg	0.89	0.90	0.89	824

The oversampling clearly increases the overfitting of the model. The F1 score and the AUC dropped more than 20% in testing. Additionally, the performance obtained by oversampling the minority class doesn't show significant benefits in testing performance than the first approach without oversampling.

Random Forrest

```
[246]: # summarize class distribution
counter = Counter(y)
print(counter)
#split_training and test set
X_train, X_test, y_train, y_test = train_test_split(training_set, y, test_size=0.
→2, random_state=42)
#compromise oversampling
over = SMOTE(sampling_strategy=0.5)
X_train, y_train = over.fit_resample(X_train, y_train)
# summarize the new class distribution
counter = Counter(y_train)
print(counter)
```

```
Counter({0: 3668, 1: 451})
Counter({0: 2936, 1: 1468})
```

```
[247]: #RF
random_forest_clf = RandomForestClassifier()

#Tuning and finding the best hyperparams
param_grid = {
    'max_leaf_nodes' : [10, 20, 30],
    'max_depth': [10, 20, 30],
}

#grid search
grid_search_rf = GridSearchCV(estimator = random_forest_clf, param_grid = _
    ↪param_grid, cv=5, scoring='f1', verbose=0, n_jobs=-1)

#training
clf_rf = grid_search_rf.fit(X_train, y_train)
y_pred_train = clf_rf.predict(X_train)
roc=roc_auc_score(y_train, y_pred_train)
print('roc', roc)
print(classification_report(y_train, y_pred_train))
```

```
roc 0.9056539509536784
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	2936
1	0.90	0.86	0.88	1468
accuracy			0.92	4404
macro avg	0.91	0.91	0.91	4404
weighted avg	0.92	0.92	0.92	4404

```
[248]: #test
y_test_pred = clf_rf.predict(X_test)

#performance
roc=roc_auc_score(y_test, y_test_pred)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_test, y_test_pred))
print(classification_report(y_test, y_test_pred))
```

```
roc 0.7389819434545023
prediction_matrix
[[692  40]
 [ 43  49]]
```

	precision	recall	f1-score	support
0	0.94	0.95	0.94	732

	1	0.55	0.53	0.54	92
accuracy				0.90	824
macro avg		0.75	0.74	0.74	824
weighted avg		0.90	0.90	0.90	824

The RF with oversampling was revealed the best model so far. Despite the overfitting of training, the precision and recall of the minority class increase compared to the previous approach, bringing a global F1 score and AUC of 0.74. In the next section, the threshold model will be adjusted to improve the performance of the minority class.

7 Performance

In the current section the best model obtained in the previous approaches was used, namely the RF with oversampling.

```
[255]: # summarize class distribution
counter = Counter(y)
print(counter)
#split_training and test set
X_train, X_test, y_train, y_test = train_test_split(training_set, y, test_size=0.
→2, random_state=42)
#compromise oversampling
over = SMOTE()
X_train, y_train = over.fit_resample(X_train, y_train)
# summarize the new class distribution
counter = Counter(y_train)
print(counter)
```

```
Counter({0: 3668, 1: 451})
Counter({0: 2936, 1: 2936})
```

```
[261]: #RF
random_forest_clf = RandomForestClassifier()

#Tuning and finding the best hyperparams
param_grid = {
    'max_leaf_nodes' : [10, 20, 30],
    'max_depth': [10, 20, 30],
}

#grid search
grid_search_rf = GridSearchCV(estimator = random_forest_clf, param_grid =_
→param_grid, cv=5, scoring='f1', verbose=0, n_jobs=-1)

#training
```

```

clf = grid_search_rf.fit(X_train, y_train)
y_train_probs = clf.predict_proba(X_train)
y_training_pred = clf.predict(X_train)

```

```

[262]: #test
y_test_pred = clf.predict(X_test)
y_test_probs = clf.predict_proba(X_test)

#set threshold evaluation
threshold = np.arange(0.35, 0.70, 0.05)
for x in threshold:
    y_pred = (clf.predict_proba(X_test)[: ,1] >= x).astype('int')
    print('test', "\n", x, "\n", classification_report(y_test, y_pred))

```

```

[272]: #performance
y_pred = (clf.predict_proba(X_test)[: ,1] >= 0.5).astype('int')

```

```

[273]: roc=roc_auc_score(y_test, y_pred)
print('roc', roc)
print('prediction_matrix', '\n', confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

```

roc 0.7810049893086244

prediction_matrix

```
[[666  66]
```

```
[ 32  60]]
```

	precision	recall	f1-score	support
0	0.95	0.91	0.93	732
1	0.48	0.65	0.55	92
accuracy			0.88	824
macro avg	0.72	0.78	0.74	824
weighted avg	0.90	0.88	0.89	824

8 Discussion

The final model was revealed to be the best model applied in the analysis, as demonstrated by the gold standard adopted for unbalanced set, such as AUC and F1 score. The AUC provides an aggregate measure of performance across all possible classification thresholds. The value of 0.78 suggests a good compromise between the True and False positive rates, with a global accuracy of 88%. In addition, the F1 score of 0.74 indicates low false positives and low false negatives. However, in the deployment phase of the model, it is still possible to adjust the threshold of the model according to the use case and business needs. In other words, if we are interested in maximizing the level of precision of the minority class, we would increase the level of the threshold of the model. In a real-world scenario, for example, if we want to be totally sure about

the client that will subscribed a term deposit, we will adopt a higher level of precision of the minority class, sacrificing its recall and global F1 score. To conclude, the final implementation of the model will be driven by the business and company needs.


9 Save Model

```
[ ]: import pickle
      filename = 'finalized_model.sav'
      pickle.dump(clf, open(filename, 'wb'))
```

10 Appendix

```
[130]: #correlation matrix of the features (to identify potential correlations among
        →them)
        corr = training_set.corr()
        # Generate a mask for the upper triangle
        mask = np.triu(np.ones_like(corr, dtype=bool))
        # Set up the matplotlib figure
        f, ax = plt.subplots(figsize=(11, 9))
        # Generate a custom diverging colormap
        cmap = sns.diverging_palette(230, 20, as_cmap=True)
        # Draw the heatmap with the mask and correct aspect ratio
        10
        sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
        square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

```
[130]: <AxesSubplot:>
```

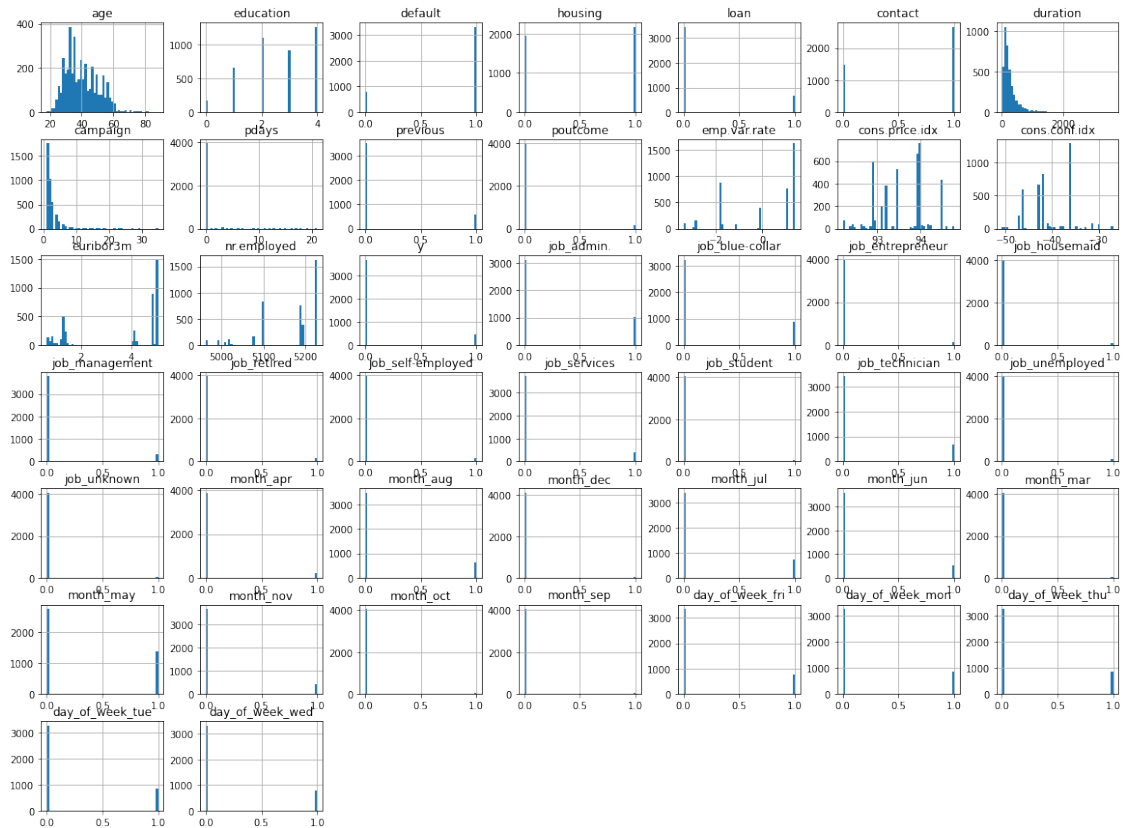


```
[146]: #checking miss values
df.isnull().sum() #no miss values detected
```

```
[146]: age      0
marital    0
education  0
default    0
housing    0
loan       0
contact    0
duration   0
campaign   0
pdays     0
previous   0
poutcome   0
emp.var.rate  0
```

```
cons.price.idx      0
cons.conf.idx       0
euribor3m           0
nr.employed         0
y                   0
job_admin.          0
job_blue-collar     0
job_entrepreneur    0
job_housemaid       0
job_management      0
job_retired         0
job_self-employed   0
job_services        0
job_student         0
job_technician      0
job_unemployed      0
job_unknown         0
month_apr           0
month_aug           0
month_dec           0
month_jul           0
month_jun           0
month_mar           0
month_may           0
month_nov           0
month_oct           0
month_sep           0
day_of_week_fri     0
day_of_week_mon     0
day_of_week_thu     0
day_of_week_tue     0
day_of_week_wed     0
dtype: int64
```

```
[147]: #checking distribution
df.hist(bins=50, figsize=(20, 15))
plt.show()
```

```
[286]: #PLOT ROC
false_positive_rate1, true_positive_rate1, threshold1 = roc_curve(y_test,
→y_test_probs[:,1])
plt.subplots(1, figsize=(10,10))
plt.title('Receiver Operating Characteristic - RandomForrest')
plt.plot(false_positive_rate1, true_positive_rate1)
plt.plot([0, 1], ls="--")
plt.plot([0, 0], [1, 0] , c=".5"), plt.plot([1, 1] , c=".5")
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

